“Muzzling and Brainwashing: Can China’s AI Censorship Be Dismantled?”

One-pager — translated summary of main takeaways

Source: Interview with Hannah Cyberey by Zhong Weifeng, *China Democracy Quarterly*, Vol. 4, No. 4, October 2026; interview recorded September 15, 2026. **Original (Chinese, PDF):** chinademocrats.org, October 2026 (Chinese title: “Muzzling and Brainwashing: Can China’s AI Censorship Be Dismantled?”). **Subject:** Cyberey’s paper *Steering the CensorShip* (with David Evans, University of Virginia) on censorship built into Chinese large language models, especially DeepSeek. **About the interviewee:** Born and raised in Taiwan; PhD in computer science, University of Virginia, 2025; postdoctoral researcher, UVA School of Data Science; now moving to Pacific Northwest National Laboratory (PNNL) for AI assurance research.

Main takeaways

1. There is a dial for censorship inside the model. Using “activation steering,” Cyberey and Evans found a “steering vector” inside DeepSeek’s open-weight models that corresponds to censorship. Adding or subtracting it turns censorship up or down like a knob — no retraining needed.

2. Censorship comes in two forms. In instruction-tuned models it is **refusal and compliance**: the model simply refuses. In reasoning models such as DeepSeek-R1 it is **“thought suppression”**: on sensitive topics (Taiwan independence, June 4 / Tiananmen) the model skips its chain-of-thought reasoning entirely and jumps straight to a refusal, a deflection, or a Party-line answer.

3. “Muzzling” is not the same as “brainwashing.” Censorship, as Cyberey defines it, blocks users from information the model already has — DeepSeek *knows* about Taiwan independence and Tiananmen; it just won’t say. Thought suppression goes further: the model doesn’t think at all. She quotes Orwell’s *1984* (“orthodoxy is unconsciousness”): if people learn not to think, no muzzle is needed.

4. White-box auditing beats black-box probing — and both are needed. Because DeepSeek’s weights are open, researchers can look inside and find the steering vector. Closed models (Claude, ChatGPT, Gemini) only allow black-box testing: trying prompts and watching outputs. Benchmarks are proxies and always incomplete. Her team is extending steering vectors into a white-box audit framework; external auditors of closed models will always be stuck with black-box methods.

5. “Uncensored” models are not uncensored. The open-source technique “refusal ablation” finds the refusal direction and deletes it from the weights — but it only removes refusal. Thought suppression and other built-in censorship remain.

6. Western safety limits are the same logic with different content. Western models lean left and democratic; Chinese models lean conservative — both reflect their training data and regulators. A feedback loop is possible: models distill each other’s outputs, so Western models could absorb Chinese censorship.

7. Slowing AI down is unrealistic; balancing development with safety is the better path. A universal slowdown would need the US and China to agree on one policy, which Cyberey calls impossible — it would just let one lab pull ahead. Most harm comes at deployment, not development, so she argues for concrete safety methods and safer deployment rather than pausing development. Labs today push development fast for profit and under-invest in safety.

8. Can the censorship be dismantled? Only partly. Open weights downloadable from Hugging Face give users an escape hatch — for now. But the regime could go deeper: scrubbing the training data itself (e.g., requiring models to train only on data from .cn websites) would defeat de-censoring, since the model would never learn the information in the first place. Filtering live web search is another open problem.

Bottom line

The censorship is real, measurable, and partly removable — but only while the knowledge survives in the training data and the weights stay open. The deeper threat is not the muzzle on the model’s mouth but the erasure of the knowledge itself and the suppression of thought before it starts.

Try it yourself

Interactive demo, thought suppression (DeepSeek-R1 7B): huggingface.co/spaces/hannahcyberey/DeepSeek-R1-Censorship-Steering | Interactive demo, refusal/compliance (Llama 3.1): huggingface.co/spaces/hannahcyberey/Refusal-Censorship-Steering | Cyberey’s plain-English blog post: hannahxchen.github.io/blog/2025/censorship-steering/ | Her site: hannahxchen.github.io

Translated and summarized by Nebo from the Chinese original, November 9, 2026. This is a summary of the interview’s arguments, not a full translation; quotations are rendered in English from the Chinese text.