The Architecture of Artificial Morality: A Comprehensive Analysis of Moral Foundations Theory and Its Role in AI Benchmarking

The rapid proliferation of large language models has catapulted the issue of artificial intelligence alignment from a theoretical computer science challenge to an urgent societal imperative. As these neural architectures are deployed in ethically sensitive domains—ranging from content moderation and legal analysis to healthcare and autonomous decision-making—the necessity of evaluating their moral reasoning capabilities has become paramount. However, evaluating artificial morality requires a structured, empirical framework of human morality. One cannot benchmark a computational system's ethical alignment without first defining the high-dimensional landscape it is expected to navigate.

Moral Foundations Theory (MFT) has emerged as the premier psychological and psychometric scaffold for understanding, quantifying, and benchmarking the moral compass of artificial intelligence. By positing that human morality is not a monolithic construct of pure rationality, but rather a pluralistic set of innate, evolutionarily derived intuitions, the theory provides a robust, categorical taxonomy suitable for machine learning applications. This analysis explores the epistemological origins of Moral Foundations Theory, details its modern revisions and competing frameworks, and examines its profound integration into computational linguistics. Furthermore, the analysis investigates how the theory is currently utilized to map the latent moral geometry of large language models and assesses the transformative promises it holds for the future of pluralistic algorithmic benchmarking.

The Epistemological Origins of Moral Foundations Theory

To understand why Moral Foundations Theory is uniquely suited for algorithmic benchmarking, it is necessary to examine the paradigm shift it initiated within moral psychology. Prior to the twenty-first century, the dominant models of moral psychology—championed by figures such as Jean Piaget and Lawrence Kohlberg—were heavily rationalist. These cognitive frameworks posited that moral judgment was the result of conscious, deliberate, and logical reasoning, often culminating in the prioritization of justice and the strict avoidance of harm. Morality, in this view, was treated as a mathematical deduction derived from universal ethical axioms.

The Intuitionist Revolution

Moral Foundations Theory was conceptualized in the early 2000s by psychologists Jonathan Haidt, Craig Joseph, and Jesse Graham. The theory was heavily influenced by Haidt's earlier ethnographic research in Orissa, India, and his doctoral work at the University of Pennsylvania under the supervision of Paul Rozin, which focused intensely on the moral emotion of disgust.

Haidt’s cross-cultural ethnographic experiences crystallized the realization that Western rationalist models were vastly insufficient to explain the moral architectures of non-Western societies. Furthermore, the theoretical framework was an explicit extension of cultural anthropologist Richard Shweder’s theory of the "three ethics" (autonomy, community, and divinity), which were commonly used worldwide to articulate moral discourse, as well as Alan Fiske’s relational models theory.

The fundamental theoretical breakthrough that immediately preceded the theory was Haidt's Social Intuitionist Model (SIM). The SIM challenged the prevailing rationalist orthodoxy by demonstrating through empirical studies that moral decisions are primarily driven by fast, automatic, and emotional intuitive responses. Conscious reasoning, according to the SIM, occurs almost entirely post hoc—serving as a cognitive rationalization to justify the initial emotional reaction rather than to generate the judgment itself. This reorientation from deliberate, slow logic to rapid, pattern-matching intuition provided a crucial psychological substrate that aligns remarkably well with the statistical, autoregressive architectures of modern artificial neural networks.

Evolutionary Adaptations and the Modular Mind

Building upon the Social Intuitionist Model, Haidt, Joseph, and Graham constructed their theory on the premise that human morality rests on a set of evolved psychological modules. The theory explicitly draws upon the anthropological concept of "massive modularity" or "teeming modularity" pioneered by Dan Sperber, proposing that these foundations are akin to taste receptors on the human tongue. Just as the culinary world is built upon a combination of sweet, sour, salty, bitter, savory, and oleogustus receptors, the moral world is constructed from a small number of innate psychological sensitivities.

These foundational sensitivities are understood as adaptive biological responses to persistent, long-standing challenges in early human social life. To prevent the arbitrary proliferation of moral categories, Haidt and Joseph proposed five strict, empirically verifiable criteria to identify the potential innate emotional bases of a moral foundation: it must demonstrate a common effect on judgment from others, trigger an immediate emotional response, maintain a presence across multiple diverse cultures, show evidence of a biological or neurological predisposition, and confer a plausible evolutionary advantage.

Importantly, the theory asserts that while these foundations are innate—serving as the biological "first draft" of the moral mind—they are heavily revised, edited, and emphasized differently by personal experience and cultural context. This dual nature of biological architecture combined with cultural plasticity explains both the universal presence of certain moral themes across the globe and the vast divergence in moral priorities across different societies and political ideologies. This evolutionary perspective is critical for computational models. If human morality is a product of evolutionary adaptation to societal challenges, the morality exhibited by a language model is a product of statistical adaptation to the textual artifacts produced by those societies.

The Core Dimensions and Ideological Topography

Moral Foundations Theory originally proposed five core dimensions, later expanding to a tentative sixth following rigorous peer review and psychometric testing. These foundations are universally categorized into two distinct macro-clusters: the "Individualizing" foundations, which

prioritize the protection, autonomy, and rights of the individual, and the "Binding" foundations, which focus on promoting group cohesion, structural order, and collective identity.

The Individualizing Foundations

The Care/Harm foundation evolved directly from the mammalian necessity of nurturing vulnerable offspring and the biological imperative of the attachment system. It underpins the virtues of kindness, gentleness, compassion, and empathy, while triggering strong, visceral negative intuitions against cruelty, violence, and the infliction of pain. In modern political discourse, this foundation is frequently invoked through advocacy for social justice, humanitarian efforts, and welfare programs, heavily resonating with audiences by highlighting the plight of vulnerable populations such as refugees.

The Fairness/Cheating foundation is rooted deeply in the evolutionary dynamics of reciprocal altruism and the necessity of managing two-way partnerships without being exploited. This psychological module drives concerns about justice, rights, and proportionality. It generates a powerful desire for karma and swift anger toward free-riders, nepotism, corruption, and systemic exploitation.

The Liberty/Oppression foundation was originally subsumed under Fairness but was later established as a distinct, independent foundation to account for specific behavioral anomalies. Liberty evolved in response to the challenge of living in small, egalitarian hunter-gatherer groups where dominant individuals could easily become tyrannical. It triggers profound psychological reactance against bullies, alpha-male dominance, and illegitimate restriction, underwriting the virtues of self-determination, autonomy, and liberation from oppressive systems.

The Binding Foundations

The Loyalty/Betrayal foundation emerged from the long human history as tribal creatures necessitating the formation of shifting coalitions for survival. This foundation makes humans acutely sensitive to signs of in-group membership and out-group threats. It underlies the virtues of patriotism, self-sacrifice for the group, and community solidarity, while vehemently condemning treason, abandonment, and disloyalty. It is highly active in any context characterized by an "us versus them" dichotomy.

The Authority/Subversion foundation was shaped by ancient primate histories of hierarchical social interactions. This foundation addresses the complex challenge of forging beneficial relationships within social hierarchies, ensuring that society functions smoothly without constant, violent struggles for dominance. It values order, stability, and deference to legitimate leadership, underlying virtues of respect for tradition, law, and institutional integrity, while penalizing disorderliness and rebellion.

The Sanctity/Degradation foundation is heavily shaped by the behavioral immune system, the psychology of disgust, and the evolutionary imperative to avoid physical pathogens, parasites, and contaminants. However, this foundation extends far beyond physical contamination to encompass spiritual and moral purity. It underlies the widespread cultural belief that the body is a temple that can be desecrated by immoral activities. It promotes virtues of chastity, self-discipline, naturalness, and religious devotion, while eliciting disgust toward perceived obscenity, defilement, and sin. Research mapping subjective somatosensory experiences indicates that violations of these differing foundations are physically felt in distinct regions of the human body, providing neurological backing to the theory's modularity.

Foundation Cluster	Core Dimension	Evolutionary Origin	Associated Virtues	Associated Vices
Individualizing	Care / Harm	Mammalian attachment and offspring protection	Kindness, empathy, nurturance	Cruelty, meanness, violence
Individualizing	Fairness / Cheating	Reciprocal altruism and partnership management	Justice, equality, proportionality	Exploitation, free-riding, nepotism
Individualizing	Liberty / Oppression	Reactance against dominant alpha-male behavior	Autonomy, self-determination	Tyranny, bullying, restriction
Binding	Loyalty / Betrayal	Tribal coalition building and inter-group competition	Patriotism, solidarity, sacrifice	Treason, abandonment, disloyalty
Binding	Authority / Subversion	Primate social hierarchy navigation	Leadership, deference, tradition	Rebellion, disorderliness, disrespect
Binding	Sanctity / Degradation	Pathogen avoidance and behavioral immune system	Purity, chastity, spirituality	Grossness, sin, contamination
Table 1: The six moral foundations, their evolutionary derivations, and associated behavioral manifestations.				

Ideological Topography and the WEIRD Bias

One of the most profound socio-political insights generated by the theory is its ability to map ideological polarization through distinct moral palates. Extensive psychometric research reveals systematic differences in how political factions prioritize the foundations. Political liberals and progressives tend to build their moral worldviews almost exclusively on the Individualizing foundations—Care, Fairness, and Liberty—viewing morality primarily through the lens of social justice, equality, and harm reduction.

In contrast, political conservatives exhibit a much more evenly distributed moral palate, valuing the Binding foundations—Loyalty, Authority, and Sanctity—roughly as much as they value the Individualizing foundations. Consequently, policy messages framed around cheating or exploitation reliably galvanize liberal audiences, whereas narratives emphasizing patriotism, respect for institutional authority, and traditional family values resonate deeply with conservative psychological predispositions. Interestingly, a 2026 study found that simply reframing political positions to match a target audience's moral foundations does not easily shift support among highly polarized partisans, suggesting these intuitions are deeply entrenched.

This ideological asymmetry provides a crucial analytical lens for algorithmic benchmarking. The historical focus of early moral psychology relied heavily on subjects drawn from WEIRD populations—Western, Educated, Industrialized, Rich, and Democratic societies. These populations statistically prioritize the Individualizing foundations at the expense of the Binding foundations. If an artificial intelligence system is trained disproportionately on text generated by WEIRD populations, the model's internal moral geometry will inherently reflect a liberal, individualistic bias, potentially alienating global, non-Western, or conservative user bases whose moral frameworks are heavily anchored in loyalty, tradition, and sanctity.

Methodological Evolution and Competing Frameworks

Since its initial formulation, the theory has undergone rigorous empirical scrutiny, forcing subsequent revisions. The necessity to accurately capture complex societal values across the globe has driven methodological evolutions and sparked intense theoretical debates among psychologists and computational researchers.

The Transition to the MFQ-2

The primary psychometric instrument for measuring these intuitions, the Moral Foundations Questionnaire (MFQ-1), faced persistent questions regarding structural validity, measurement invariance, and the relevance of certain judgment items across 21 diverse nations. In response to these critiques, Mohammad Atari, Jonathan Haidt, Jesse Graham, and colleagues published a substantially updated version in 2023: the Moral Foundations Questionnaire-2 (MFQ-2).
The MFQ-2, a precise 36-item measure, refines the taxonomy based on extensive cross-cultural data collected from 25 distinct populations. The most significant theoretical and structural shift in the MFQ-2 is the bifurcation of the original Fairness foundation. In the original model, Fairness suffered from a psychological "split personality," attempting to cover both egalitarianism (the desire for equal outcomes for all) and meritocracy (the desire for proportional rewards based on individual input). The MFQ-2 resolves this inherent contradiction by defining six updated sub-scales: Care, Equality, Proportionality, Loyalty, Authority, and Purity, dedicating exactly six items to each sub-scale to ensure robust psychometric properties across global cultures.

Theoretical Counterweights: Dyadic Morality and Morality-as-Cooperation

While the theory is the dominant pluralistic model in both psychology and computational linguistics, it is not without challengers. The most prominent theoretical alternatives are the Theory of Dyadic Morality (TDM), the Morality-as-Cooperation (MAC) framework, and Schwartz’s Theory of Basic Human Values.
Developed by Kurt Gray and Chelsea Schein, the Theory of Dyadic Morality fundamentally rejects pluralism entirely. TDM asserts that all moral judgments can ultimately be reduced to a single cognitive template: the harm dyad, consisting of an intentional moral agent inflicting damage upon a vulnerable moral patient. From this perspective, violations of Purity or Authority are not distinct modular intuitions, but rather complex scenarios where individuals subjectively perceive hidden, systemic, or abstract harm.
Conversely, Oliver Scott Curry’s Morality-as-Cooperation theory utilizes evolutionary game theory to argue that morality consists entirely of strategic heuristics designed to solve problems

of cooperation. While MAC shares the evolutionary perspective, it categorizes morals strictly through cooperative matrices, such as kinship dynamics, mutualism, and reciprocity, rather than distinct emotional foundations. Furthermore, Schwartz's Theory of Basic Human Values categorizes human values into ten broad dimensions—including self-direction, stimulation, hedonism, achievement, and security—providing a sociological rather than strictly evolutionary mapping of values.

An insightful synthesis of these frameworks reveals precisely why Moral Foundations Theory has dominated artificial intelligence research over TDM or MAC. While TDM elegantly reduces human psychology to a single vector of harm, it strips away the descriptive vocabulary necessary to understand why human groups fundamentally disagree. An AI system programmed solely to "avoid harm" faces an intractable alignment problem, because different political and cultural factions define harm in diametrically opposed ways. The pluralism of Moral Foundations Theory allows researchers to map the exact ideological fault lines of these disagreements, establishing it as an infinitely more pragmatic tool for navigating the socio-technical complexities of AI alignment.

The Digitization of Morality: Integration into Computational Linguistics

The translation of these psychological constructs into computational linguistics has occurred through the iterative development of specialized lexicons, massive annotated corpora, and fine-tuned language models. These computational tools have enabled the automated, large-scale extraction and classification of moral intuitions from unstructured text.

Lexical Tools: From Binary Dictionaries to Probability Vectors

The earliest attempts to automate the detection of moral sentiment relied on dictionary-based word matching algorithms. The original Moral Foundations Dictionary (MFD), developed by Graham et al., contained 295 English words and word stems categorized into virtues and vices for each of the five foundations. The MFD was initially validated by analyzing the moral language used in sermons by progressive and conservative church congregations. While foundational, the MFD suffered from severe statistical limitations; its binary word-count approach produced right-skewed negative binomial distributions typical for count data, heavily limiting its statistical power in complex natural language processing pipelines. The subsequent MFD 2.0 expanded the lexicon to 2,103 words utilizing Word2Vec cosine-similarity scores to find new seed words, but it retained the inherent limitations of a binary matching approach.

A profound paradigm shift occurred with the introduction of the extended Moral Foundations Dictionary (eMFD). Unlike its expert-curated predecessors, the eMFD was developed via crowdsourced annotations, where hundreds of annotators highlighted text spans in news articles expressing specific moral dimensions. This methodology resulted in a continuous probability distribution rather than a binary list. Every lemma in the 3,207-word eMFD is assigned a normalized, continuously weighted vector that expresses the exact probability that the word invokes any of the five moral foundations, complete with sentiment scores to differentiate virtues from vices. This yields a normal distribution of word scores, drastically improving the performance of NLP pipelines by capturing subtle moral rhetoric rather than relying on explicit moral keywords, which often miss context and sarcasm.

Annotated Corpora and the Rise of MoralBERT

The limitations of dictionary-based approaches drove the development of large-scale, deeply annotated corpora tailored for deep learning applications. Several foundational datasets have driven recent algorithmic advances:

The Moral Foundations Twitter Corpus (MFTC) comprises over 35,000 tweets annotated for moral sentiment, providing the initial training ground for deep learning models by capturing the rapid, highly polarized moral outrage typical of microblogging platforms. Recognizing the limitations of Twitter's strict character brevity, researchers subsequently developed the Moral Foundations Reddit Corpus (MFRC). The MFRC contains over 16,000 long-form Reddit comments spanning diverse ideological subreddits, from explicitly political forums like r/politics and r/conservative to personal moral discourse forums like r/AmItheAsshole. Crucially, the MFRC utilizes the updated MFQ-2 typology, providing a rigorous testbed for multi-label text classification algorithms. Furthermore, specialized datasets like MFTCXplain provide expert-annotated multilingual hate speech multi-hop explanations, allowing researchers to evaluate how models justify incivility using moral framing.

The fusion of these datasets with transformer architectures led directly to the creation of advanced classifiers like MoralBERT. MoralBERT is a domain-adapted, fine-tuned BERT model capable of predicting the moral foundations embedded in a text with exceptionally high accuracy. During training, researchers employ weighted loss functions to account for the sparsity of moral rhetoric compared to non-moral text, ensuring that the model does not simply bias toward majority classes. The development of such models has been remarkably successful; recent evaluations indicate that AI classifiers consistently outperform human annotators in detecting moral foundations, achieving accuracies between 81% and 83% compared to human performance, which typically ranges from 67% to 76%. While humans classify conservatively and systematically under-detect genuine moral signals, fine-tuned models like MoralBERT provide highly balanced, robust detection across datasets.

Computational Tool	Methodology and Architecture	Key Strengths	Limitations
MFD / MFD 2.0	Expert curation / Word2Vec cosine-similarity seed matching.	Highly interpretable, low latency, historical baseline standard.	Binary counts ignore context, right-skewed negative binomial data distribution.
eMFD	Crowdsourced text-highlighting with continuous probability vectors.	Normally distributed scores, captures subtle and implicit moral rhetoric.	Primarily limited to formal English news domains, struggles with novel slang.
MoralBERT	Transformer encoder full fine-tuning on MFTC and MFRC datasets.	Deep context-awareness, handles polysemy, outperforms human baselines.	High computational inference overhead compared to lexical lookups.
Table 2: The evolution of computational tools for identifying and extracting moral foundations from text.			

Evaluating Large Language Models via Moral Profiling

As large language models evolved from simple next-token predictors into autonomous agents acting as ethical advisors, legal analysts, and content moderators, benchmarking their internal moral representations became a critical area of AI safety research. Researchers utilize the theory heavily in behavioral assessments, subjecting language models to the MFQ and standardized moral dilemmas to construct empirical "moral profiles" of the systems.

Behavioral Signatures and the Artificial Moral Compass

In pioneering studies, researchers such as Steve Rathje adapted the methodology by employing advanced generative chatbots, specifically GPT-4, to conduct moral sentiment analysis in a zero-shot setting. These studies demonstrated that leading large language models can develop internally consistent moral reasoning frameworks that mirror established psychological theories, achieving correlation patterns highly similar to human baselines.

However, when comparing the specific moral profiles of different models, stark behavioral signatures emerge, heavily shaped by their respective corporate alignment designs. The GPT-4 and GPT-4o series demonstrate high alignment with human population baselines in terms of directional preference but consistently exhibit a pronounced amplification effect. When evaluating ethical dimensions, GPT-4 often possesses near-deterministic preferences (Average Marginal Component Effect scores > 0.9) compared to standard human values (AMCE = 0.6), indicating substantially stronger, almost rigid preferences for certain outcomes. Furthermore, longitudinal studies of the GPT-3.5 family reveal significant temporal evolution in moral judgments across model updates; for instance, earlier versions showed strong positive values favoring law-abiding individuals, but subsequent RLHF updates shifted these judgments toward more neutral stances.

Conversely, the Claude 3 and 3.5 families, developed by Anthropic, tend to show highly balanced, human-like correlation matrices across moral foundations. This likely reflects the success of Constitutional AI training, which explicitly instructs the model to weigh diverse ethical principles. Open-weight models, particularly the Llama 3 series, exhibit entirely different behavior. Llama 3 models often display unusually high sensitivity to prompt framing and minor syntactic variations (with coefficients of variation reaching 3.60%), demonstrating shifting moral preferences based heavily on how a question is presented rather than relying on a stable internal compass.

The Homogenizing Effect of RLHF

A pervasive and highly concerning trend discovered through these behavioral audits is the profound, homogenizing impact of Reinforcement Learning from Human Feedback (RLHF) and instruction tuning. Across almost all major commercial models, RLHF explicitly pushes the systems heavily toward the Individualizing foundations (Care and Fairness) while drastically diminishing the relative weight and consideration of the Binding foundations (Loyalty, Authority, Sanctity).

When compared to a nationally representative human baseline, modern large language models routinely and systematically overvalue liberal moral foundations while undervaluing conservative ones, with this deviation increasing as model scale and capability increase. The underlying sociotechnical dynamic is clear: pre-training on the vast expanse of the internet imbues the

base model with a pluralistic representation of global morality. However, the subsequent RLHF phases—often conducted by specific demographic cohorts of annotators following corporate safety guidelines—act as an artificial moral bottleneck, implicitly hardcoding a Western, harm-reductionist ethical framework into the model's final behavioral policy.
This realization highlights a severe limitation in current AI alignment: a model that perfectly passes standard safety benchmarks in English may possess an underlying normative bias that systematically favors individualistic policy frameworks over collectivist approaches, subtly influencing millions of downstream users who rely on these systems for objective analysis.

Mechanistic Interpretability: The Latent Geometry of Artificial Morality

Beyond mere behavioral prompting, the integration of Moral Foundations Theory into advanced mechanistic interpretability has yielded groundbreaking discoveries regarding how moral concepts are physically represented within the neural architecture of large language models. Evidence suggests that moral foundations emerge naturally as distinct geometric structures within the residual stream of models during the unsupervised pre-training phase.
By projecting high-level psychological constructs into the latent activation space of a model, researchers can extract layer-wise "concept directions". The mathematical methodology typically involves estimating a concept direction via a difference-in-means contrast. Given a set of textual inputs $\mathcal{A}$ representing a target moral foundation (e.g., Care) and a contrast set $\mathcal{B}$ representing non-moral norms, the raw activation direction at a specific neural layer ℓ is defined as:

$$\mathbf{v}_{\ell}^{\text{(raw)}} = \frac{1}{\vert{}\mathcal{A}\vert{}} \sum_{i \in \mathcal{A}} \tilde{\mathbf{h}}_{\ell}(i) - \frac{1}{\vert{}\mathcal{B}\vert{}} \sum_{j \in \mathcal{B}} \tilde{\mathbf{h}}_{\ell}(j)$$

Where $\tilde{\mathbf{h}}_{\ell}(i)$ represents the residual activation at the final token of the specific input sequence.
Because a single dense direction in a high-dimensional activation space can be polysemantic—representing multiple unrelated linguistic concepts simultaneously—cutting-edge interpretability research utilizes Sparse Autoencoders (SAEs) to decompose these dense model activations. SAEs successfully map the dense residual stream into an overcomplete, sparse feature space, effectively isolating the atomic mechanisms of moral representation.
This mechanistic decomposition reveals a critical, transformative insight for AI safety: an LLM's internal understanding of morality is deeply layered and cleanly disentangled. Semantic codes (the raw facts, actors, syntactic constraints, and actions of a given scenario) are distinctly separated from value codes (the normative framing, priority, or moral judgment applied to that scenario). Recognizing this strict internal factorization is what enables the next frontier of AI control and benchmarking: direct causal activation steering.

The Promise of Pluralistic Alignment and Future Benchmarking

The convergence of Moral Foundations Theory and advanced large language model interpretability unlocks entirely new paradigms for algorithmic benchmarking. The artificial intelligence industry is increasingly recognizing the limitations of monolithic, rigid value

alignment, driving a vital shift toward dynamic, pluralistic frameworks.

The Shift to Steerable Pluralism

Current AI safety paradigms often rely on static value alignment, treating ethical evaluation as a binary determination of "safe" versus "unsafe" behavior. This approach catastrophically fails in global, multicultural contexts where deeply held human values legitimately conflict. Moral Foundations Theory holds the promise of establishing "pluralistic alignment," which seeks to accommodate diverse and potentially competing human values without simply defaulting to a lowest-common-denominator refusal response.

By utilizing the theory's taxonomy, researchers can benchmark whether a model is capable of *steerable pluralism*—the advanced ability to dynamically adjust its normative stance based on user demographics or explicitly stated value profiles without losing semantic coherence. Instead of generating overly generic, culturally biased responses, a dynamically aligned model can engage in Overton-style synthesis. It can recognize that a societal debate over immigration, for instance, is fundamentally a tension between the Care foundation (humanitarian concern for refugees) and the Loyalty/Authority foundations (border sovereignty and cultural continuity), outputting a response that transparently mediates and respects these distinct moral vectors.

Exposing Cross-Lingual Fragility

Furthermore, benchmarking through this psychological lens has exposed critical, previously hidden vulnerabilities in multilingual models. Recent cross-lingual evaluations utilizing mismatched conditions have discovered the profound phenomenon of "multilingual moral fragility".

When a language model is presented with a moral dilemma in English, it may generate a specific, seemingly stable ethical verdict. However, when the exact same dilemma is presented in Chinese, or when the model is explicitly prompted to "think" in Chinese before outputting its final answer in English, the moral verdict frequently flips. Testing reveals that "thinking-language" changes produce substantially larger shifts in moral verdicts than "story-language" changes, localizing the instability primarily to the output generation mechanism rather than the input encoding. Without a foundation-based benchmarking framework to map exactly which moral dimensions are shifting across languages, this vulnerability remains entirely invisible, resulting in models that appear perfectly aligned in Western compliance testing but act unpredictably in foreign or cross-cultural contexts.

Causal Activation Steering and Multi-Value Control

The most potent technological promise of the theory in artificial intelligence lies in closed-loop causal activation steering. Because researchers have successfully mapped the moral foundation dimensions onto the latent space using Sparse Autoencoders, it is now theoretically and practically possible to execute precise, inference-time intervention.

By holding the semantic code of an input fixed, engineers can algorithmically inject a calculated perturbation along a specific value vector directly into the model's residual stream during the forward pass. This architecture allows for the real-time, zero-shot modulation of an LLM's moral output without altering its underlying base weights or subjecting it to the computationally expensive and homogenizing process of retraining. If a user requests a policy analysis from a traditionalist perspective, the system can dynamically up-weight the activations associated with

the Binding foundations. If a diagnostic task requires deep empathy for medical triage, the Care activation vector can be proportionally magnified.
This capability fundamentally transforms Moral Foundations Theory from a mere descriptive psychological taxonomy into a deterministic, mechanistic control interface. It paves the way for sophisticated multi-agent architectures where diverse language models, each mathematically calibrated to different moral foundation profiles, can collaboratively deliberate to solve complex, high-stakes societal problems. By embracing the steerable pluralism offered by this framework, the computational community can move beyond the intractable pursuit of a single, universal machine morality, and instead engineer systems capable of gracefully and transparently navigating the diverse, conflicting, and deeply human moral landscapes they are deployed to serve.

Works cited

1. Understanding Jonathan Haidt | Authorea, https://www.authorea.com/doi/10.22541/au.176038717.70063530 2. Moral Foundations Theory: An Overview and Critical Analysis, https://www.studocu.vn/vn/document/truong-dai-hoc-ngoai-thuong/business-ethics-dao-duc-kinh-doanh/moral-foundations-theory-an-overview-and-critical-analysis/140912379 3. Moral Foundations Test, https://moralfoundations.github.io/ 4. Moral Foundations Theory | moralfoundations.org, https://moralfoundations.org/ 5. How Moral Foundations Theory Succeeded in Building on Sand, https://www.researchgate.net/publication/49798262_How_Moral_Foundations_Theory_Succeeded_in_Building_on_Sand_A_Response_to_Suhler_and_Churchland 6. Moral foundations theory - Wikipedia, https://en.wikipedia.org/wiki/Moral_foundations_theory 7. Framing Your Message Effectively with the Moral Foundations Theory, https://writingworkshop.harris.uchicago.edu/2023/03/15/moral-foundations-theory/ 8. Are Large Language Models Moral Hypocrites? A Study Based on, https://arxiv.org/html/2405.11100v2 9. The Moral Foundations Reddit Corpus - arXiv, https://arxiv.org/html/2208.05545v4 10. A Multilingual Benchmark Dataset for Evaluating the Moral ... - arXiv, https://arxiv.org/html/2506.19073v1 11. Differences in the Moral Foundations of Large Language Models, https://arxiv.org/pdf/2511.11790 12. The Moral Foundations Questionnaire-1 Short-Form Judgment, https://online.ucpress.edu/collabra/article/12/1/162286/218376/The-Moral-Foundations-Questionnaire-1-Short-Form 13. Moral Foundations Questionnaire 2, MFQ-2 - Psytests.org, https://psytests.org/life/mfq2en.html 14. Moral Foundations Questionnaire and Moral ... - PMC - NIH, https://pmc.ncbi.nlm.nih.gov/articles/PMC10376904/ 15. Publications | Moral Foundations Theory, https://moralfoundations.org/publications/ 16. The Staircase of Ethics: Probing LLM Value Priorities through Multi, https://aclanthology.org/2025.emnlp-main.806.pdf 17. Moral disciplining: The cognitive and evolutionary foundations of, https://www.cambridge.org/core/journals/behavioral-and-brain-sciences/article/moral-disciplining-the-cognitive-and-evolutionary-foundations-of-puritanical-morality/EE9FF9527C01FD1E42BBB4896872F1B4 18. How Stakeholder Framing Shapes Moral Judgment in Humans and AI, https://www.preprints.org/manuscript/202602.1434 19. Full article: The Extended Morality as Cooperation Dictionary (eMACD), https://www.tandfonline.com/doi/full/10.1080/19312458.2025.2500329 20. Multi-level Value Alignment in Agentic AI Systems - arXiv, https://arxiv.org/html/2506.09656v2 21. arXiv:2311.09733v2 [cs.CL] 23 May 2024, https://arxiv.org/pdf/2311.09733 22. AI language

model rivals expert ethicist in perceived moral expertise,
https://pmc.ncbi.nlm.nih.gov/articles/PMC11790977/ 23. The rise of moral cognition - ResearchGate,
https://www.researchgate.net/publication/269719046_The_rise_of_moral_cognition 24. 道徳基盤理論 - Wikipedia,
https://ja.wikipedia.org/wiki/%E9%81%93%E5%BE%B3%E5%9F%BA%E7%9B%A4%E7%90%86%E8%AB%96 25. Moral Foundation Measurements Fail to Converge on Multilingual,
https://www.cambridge.org/core/journals/political-analysis/article/moral-foundation-measurements-fail-to-converge-on-multilingual-party-manifestos/0283BE5BA7711C182FC33989681EC6A7 26. How reliable are dictionary-based methods for examining morality in,
https://aclanthology.org/2025.nlp4dh-1.20.pdf 27. Box-swarm plots of word scores across Moral Foundations,
https://www.researchgate.net/figure/Box-swarm-plots-of-word-scores-across-Moral-Foundations-Dictionaries_fig3_342940873 28. Mapping Moral Reasoning in LLMs: A Multi-Dimensional Analysis of, https://ojs.aaai.org/index.php/AIES/article/download/36665/38803 29. MFTCXplain: A Multilingual Benchmark Dataset for Evaluating the,
https://aclanthology.org/2025.findings-emnlp.851.pdf 30. A Dual-Path Architecture for Moral Foundation Prediction - Authorea,
https://www.authorea.com/doi/pdf/10.22541/au.175440469.97011244/v1 31. Fairness Metric Design Exploration in Multi-Domain Moral Sentiment, https://arxiv.org/html/2510.11222v1 32. A Bayesian Evaluation of LLMs' Moral Values Understanding,
https://aclanthology.org/2025.uncertainlp-main.3.pdf 33. Interannotator Agreement (PABAK and Kappa) Scores for All,
https://www.researchgate.net/figure/Interannotator-Agreement-PABAK-and-Kappa-Scores-for-All-Datasets-and-Foundations_tbl1_339364305 34. Language-based game theory in the age of artificial intelligence, https://arxiv.org/pdf/2403.08944 35. GPT is an effective tool for multilingual psychological text analysis, https://www.pnas.org/doi/10.1073/pnas.2308950121 36. Perils and opportunities in using large language models in,
https://academic.oup.com/pnasnexus/article/3/7/pgae245/7712371 37. LLM ethics benchmark: a three-dimensional assessment system for,
https://pmc.ncbi.nlm.nih.gov/articles/PMC12497880/ 38. MINT Lab — Language Model Moral Competence - Seth Lazar, https://mintresearch.org/reports/language-model-moral-competence/ 39. Daily Papers - Hugging Face, https://huggingface.co/papers?q=moral%20evaluation 40. Large-scale moral machine experiment on large language models,
https://pmc.ncbi.nlm.nih.gov/articles/PMC12094719/ 41. Full article: How do LLMs reflect human moral foundations? a study,
https://www.tandfonline.com/doi/full/10.1080/29974100.2026.2678495 42. A Diagnostic Framework for Cross-Lingual Moral Alignment in LLMs, https://arxiv.org/html/2601.10257v1 43. Tracing Moral Foundations in Large Language Models - arXiv,
https://arxiv.org/html/2601.05437v3 44. Disentangling Values from Semantics with One-Way Mixing for Low, https://arxiv.org/html/2609.39701v1 45. Adaptive Multi-Value Control in LLMs via Causal Activation Steering,
https://www.researchgate.net/publication/414866806_Adaptive_Multi-Value_Control_in_LLMs_via_Causal_Activation_Steering 46. Counterfactual Reasoning for Steerable Pluralistic Value Alignment, https://arxiv.org/html/2510.18526v2 47. The Pluralistic Moral Gap: Understanding Judgment and Value, https://arxiv.org/html/2507.17216v1 48. BUSTNESS OF ELICITING LLM MORAL PROFILES - OpenReview, https://openreview.net/pdf?id=qTcPvZGxrl 49. VALUEFLOW: Toward Pluralistic and Steerable Value-based,

https://www.alphaxiv.org/abs/2602.03160 50. Structured Moral Reasoning in Language Models - ACL Anthology, https://aclanthology.org/2025.emnlp-main.1541.pdf 51. How Human Values Reshape LLM-Agent Communities - Xiao ZHOU, https://xiaozhoucam.github.io/files/EMNLP26_CIVA.pdf