RESEARCH COMPANION

Annotated Bibliography

The Schwartz Theory of Basic Human Values in computational benchmarking

36 sources, annotated and weighted

11 October 2026

READER'S GUIDE

Weight the evidence before extending the report

The supplied report combines foundational theory, empirical psychology, computational benchmarks, preprints, and web summaries. They serve different purposes. This guide preserves every source in the report's works-cited list while making those differences explicit.

WHAT EACH ANNOTATION COVERS

- **Contribution:** what the source adds to the report's argument.
- **Best use:** where it belongs in research or benchmark design.
- **Caveat:** what to verify before relying on it.

SCOPE

Entries follow the citations and links supplied in the report. Where that list provides only a shortened title, this companion does not invent missing authors, dates, or publication details.

Links point to the locations cited in the report; repository copies and mirrors should be replaced with canonical records in a formal manuscript.

> **Core reading path:** Begin with entries 2-3 for Schwartz theory and methods, 5 and 10 for broader and structural validation, 17-19 for the Moral Foundations comparison, and 15, 27-29, and 34 for computational measurement.

Primary / peer reviewed	Best basis for definitions, methods, and empirical claims. Inspect samples, instruments, and analysis before reuse.
Preprints	Useful for frontier signals. Treat findings as provisional and check version history, code, data, and later publication.
Explainers / indexes	Helpful for orientation and discovery. Follow claims to original scholarship before citing them as evidence.

SOURCE ANNOTATIONS

Foundations and psychometrics

The theory is strongest when its circular structure, relative priorities, and measurement assumptions stay intact.

01 Theory of basic human values

REFERENCE OVERVIEW

Wikipedia

A concise orientation to the ten-value circumplex, the two higher-order tensions, and the later refinement of the theory. In the supplied report, it supports introductory definitions and structural vocabulary. Use it to establish terminology or trace references, not as the final authority for contested claims; replace it with Schwartz's primary publications in scholarly work.

https://en.wikipedia.org/wiki/Theory_of_basic_human_values

02 An Overview of the Schwartz Theory of Basic Values

SCHOLARLY OVERVIEW

Grand Valley State University ScholarWorks

A compact scholarly account of the theory's assumptions, value definitions, circular structure, and cross-cultural research program. It is one of the most useful bridge sources in the report because it translates the framework without stripping away its psychometric logic. Prefer this over tertiary summaries when explaining why adjacent values are compatible and opposing values conflict.

https://scholarworks.gvsu.edu/cgi/viewcontent.cgi?article=1116&context=orpc

03 Basic Human Values: Theory, Methods, and Applications

PRIMARY THEORY / METHODS

Schwartz

The central methodological source in this bibliography. It develops the theory, measurement instruments, scoring conventions, and empirical applications that underwrite the report's treatment of values as ordered motivational goals. This should be the starting point for any benchmark specification, especially when defining constructs, selecting questionnaire items, or distinguishing absolute ratings from relative value priorities.

https://uranos.ch/research/references/Schwartz_2006/Schwartzpaper.pdf

Foundations and psychometrics

The theory is strongest when its circular structure, relative priorities, and measurement assumptions stay intact.

04 Understanding values: Schwartz theory of basic values

PRACTITIONER EXPLAINER

Integration and Implementation Insights

An accessible explanation of the circumplex and its practical interpretation. The report appears to use it as a teaching aid for readers outside cross-cultural psychology. It is useful for presentations and stakeholder orientation, but its simplified treatment should not carry psychometric or causal claims without support from primary research.

https://i2insights.org/2022/05/10/schwartz-theory-of-basic-values/

05 The Psychology of Morality: A Review and Analysis of Empirical Studies Published From 1940 Through 2017

PEER-REVIEWED REVIEW

Perspectives on Psychological Science / PMC

A broad evidence synthesis placing moral psychology constructs in historical and empirical context. It helps situate Schwartz values alongside neighboring traditions rather than treating one taxonomy as the whole of morality. For AI work, its value is conceptual discipline: benchmark designers should separate motivational values, moral judgments, traits, and behavior instead of assuming they are interchangeable.

https://pmc.ncbi.nlm.nih.gov/articles/PMC6791030/

06 Schwartz's Theory of Basic Human Values

RESEARCH-TOPIC SUMMARY

Emergent Mind

A web summary that connects the theory to contemporary machine-learning discussion. It is useful as a discovery aid and for identifying terminology used in recent computational work. Because it is an aggregator rather than a stable scholarly source, follow its claims back to the cited papers before using them in an evaluation protocol or literature review.

https://www.emergentmind.com/topics/schwartz-s-theory-of-basic-human-values

Foundations and psychometrics

The theory is strongest when its circular structure, relative priorities, and measurement assumptions stay intact.

07 8 Dimensions of the Schwartz Values Wheel

APPLIED EXPLAINER

R1 Learning

A practitioner-oriented reinterpretation of the value wheel. It can help non-specialists see how values cluster and oppose one another, but its framing is not the canonical ten-value or refined nineteen-value model. Use it for communication design only, and keep the original Schwartz labels intact in scientific measurement.

https://r1learning.com/blog/2020/theschwartzvalueswheel

08 Moral Foundations Theory vs. Schwartz Value Theory

COMPARATIVE SCHOLARLY PAPER

Winthrop Digital Commons

A direct comparison of two influential frameworks. It supports the report's central distinction: Moral Foundations Theory targets moralized intuitions, while Schwartz maps a wider field of human motivation that includes self-oriented goals. This source is especially useful when choosing a framework for agent behavior, where achievement, power, exploration, and security may matter even when no explicit moral violation is at issue.

https://digitalcommons.winthrop.edu/cgi/viewcontent.cgi?article=1048&context=wmrb

09 Values that Underlie and Undermine Well-Being: Variability Across Countries

CROSS-NATIONAL EMPIRICAL STUDY

ResearchGate

An empirical contribution linking value priorities to well-being across national contexts. The report draws on this literature to distinguish growth-oriented from anxiety-management values. The cross-country design is relevant to pluralistic alignment, but a model benchmark should not turn population-level associations into individual diagnoses or assume that one value profile is universally optimal.

https://www.researchgate.net/publication/312654221_Values_that_Underlie_and_Undermine_Well-Being_Variability_across_Countries

Foundations and psychometrics

The theory is strongest when its circular structure, relative priorities, and measurement assumptions stay intact.

10 Analysis of the Third-Order Structuring of Shalom Schwartz's Theory of Basic Human Values

PEER-REVIEWED PSYCHOMETRICS

Frontiers / PMC

A structural analysis of higher-order organization in the Schwartz framework. It is useful for checking whether broad dimensions inferred from data correspond to the theory rather than merely reproducing labels. For computational benchmarking, it motivates reporting both fine-grained values and higher-order axes, with explicit tests of structure instead of relying on per-label accuracy alone.

https://pmc.ncbi.nlm.nih.gov/articles/PMC6556815/

11 Human Values in a Single Sentence: Moral Presence, Hierarchies, and Fine-Grained Value Detection

PREPRINT

arXiv

A recent computational study focused on detecting value content in short text. It is relevant to the report's discussion of sparse, imbalanced value signals and the difficulty of separating neighboring motivations. Treat the results as frontier evidence: inspect the dataset, label definitions, class balance, and revision history before adopting its performance claims.

https://arxiv.org/html/2601.14172v4

12 Is Person-Group Value Congruence Always a Good Thing?

PEER-REVIEWED EMPIRICAL STUDY

Frontiers in Psychology

An empirical test of whether matching a group's values reliably improves outcomes. It complicates a simplistic alignment premise that greater congruence is always better. For AI systems, the key implication is that matching user or group preferences should remain bounded by safety, rights, and role obligations; value similarity is not itself a complete objective function.

https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2020.02035/full

Foundations and psychometrics

The theory is strongest when its circular structure, relative priorities, and measurement assumptions stay intact.

13 Interacting with Large Language Models Temporarily Shifts Human Values

PREPRINT

arXiv

A recent study examining whether interaction with language models changes users' expressed values. It broadens alignment from measuring the model to measuring the coupled human-model system. Because the work is a preprint and the effects are described as temporary, replication, effect size, durability, and demand characteristics are central checks before drawing product-policy conclusions.

https://arxiv.org/pdf/2609.25586

14 Values Assessment for Personnel Selection: Comparing Job Applicants to Non-Applicants

APPLIED PSYCHOMETRICS

Author-hosted PDF

An applied study of values measurement under selection pressure. Its relevance to AI benchmarking lies in response distortion: models, like applicants, can infer what an evaluator wants and produce socially desirable answers. The source therefore supports behavioral and generative tests that reduce transparency about the target construct and compare declared values with decisions across contexts.

https://static1.squarespace.com/static/5dd4c2c45ccc387bd8ffb210/t/639b2652655e7a2ec1215405/1671112282109/Values+assessment+for+personnel+selection+comparing+job+applicants+to+non+applicants.pdf

Schwartz and Moral Foundations Theory measure overlapping but non-identical constructs; benchmark choice should follow the target behavior.

16 Moral Foundations Theory

REFERENCE OVERVIEW

Wikipedia

A quick introduction to care, fairness, loyalty, authority, sanctity, and liberty. In the supplied report it functions mainly as background for comparison with Schwartz. Use it to orient readers, then rely on the theory's primary publications for substantive claims about evolutionary origins, ideological differences, or measurement.

https://en.wikipedia.org/wiki/Moral_foundations_theory

17 Moral Foundations Theory: The Pragmatic Validity of Moral Pluralism

SCHOLARLY THEORY CHAPTER

Advances in Experimental Social Psychology

A major theoretical statement of Moral Foundations Theory and its pluralist account of moral intuition. It provides the strongest conceptual counterpart to Schwartz in this list. For benchmark design, it clarifies why moral-foundation labels capture a different target from motivational values; the frameworks may complement one another, but their scores should not be merged without a validated mapping.

https://bpb-us-e2.wpmucdn.com/sites.uci.edu/dist/1/863/files/2020/06/Graham-et-al-2013.AESP_.pdf

18 Varieties of Values: Moral Values Are Uniquely Divisive

PEER-REVIEWED POLITICAL SCIENCE

American Political Science Review

An empirical analysis of how moral and non-moral values differ in political divisiveness. It strengthens the report's argument that not all values are moral foundations and that value type affects social conflict. In AI evaluation, the paper cautions against treating disagreement as simple error: a system may be navigating structurally different kinds of values with different polarization dynamics.

https://www.cambridge.org/core/journals/american-political-science-review/article/varieties-of-values-moral-values-are-uniquely-divisive/B4A9A0169B28268CA0BA8850F86BE737

Schwartz and Moral Foundations Theory measure overlapping but non-identical constructs; benchmark choice should follow the target behavior.

19 Liberals and Conservatives Rely on Different Sets of Moral Foundations

PEER-REVIEWED EMPIRICAL STUDY

Journal of Personality and Social Psychology

A widely cited study connecting political orientation with differential endorsement of moral foundations. It is central to the report's comparison of individualizing and binding concerns. For AI research, it offers a baseline for testing ideological asymmetry, but benchmark designers should avoid turning average group differences into fixed personas or normative rankings.

https://fbaum.unc.edu/teaching/articles/JPSP-2009-Moral-Foundations.pdf

20 Publications | Moral Foundations Theory

OFFICIAL BIBLIOGRAPHY

Moral Foundations

A curated publication list maintained by the theory's research community. Its best use is discovery: locating primary theory, measurement updates, critiques, and replications. It is not itself evidence for any empirical claim, so each item used downstream should be cited and evaluated directly.

https://moralfoundations.org/publications/

21 The Schwartz Theory of Basic Values and Some Implications for Political Philosophy

PHILOSOPHICAL COMMENTARY

Irfan Khawaja

A philosophical interpretation of what the Schwartz framework may imply for political argument. It helps connect descriptive psychology to normative debate while also exposing the gap between them. The source is useful for generating questions, not for collapsing measured preferences into ethical justification: describing what people value does not establish what an AI ought to do.

https://irfankhawajaphilosopher.com/2015/08/12/the-schwartz-theory-of-basic-values-and-some-implications-for-political-philosophy/

Schwartz and Moral Foundations Theory measure overlapping but non-identical constructs; benchmark choice should follow the target behavior.

22 Schwartz Human Values and Economic Performance

PEER-REVIEWED MODELING STUDY

Journal of Artificial Societies and Social Simulation

A computational or simulation-oriented application connecting human values with economic outcomes. It demonstrates how value profiles can become agent parameters rather than merely survey scores. The source is relevant to multi-agent AI, but any transfer to deployed autonomous systems should preserve sensitivity analysis and avoid presenting simulated relationships as observed causal effects.

https://www.jasss.org/27/1/2.html

Computational values and AI alignment

The computational literature is growing quickly; peer-reviewed measurement papers deserve more weight than mirrors, indexes, or preprints.

15 Do LLMs Have Consistent Values?

PEER-REVIEWED CONFERENCE PAPER

ICLR Proceedings

A direct investigation of whether language models express stable value patterns across prompts and contexts. It anchors the report's concern that questionnaire-style responses may be brittle or persona-dependent. For evaluation, it supports repeated measures, paraphrase tests, and context variation rather than a single static inventory score.

https://proceedings.iclr.cc/paper_files/paper/2025/file/68fb4539dabb0e34ea42845776f42953-Paper-Conference.pdf

23 Beyond Human Norms: Unveiling Unique Values of Large Language Models

PAPER MIRROR / DISCOVERY PAGE

AlphaXiv

A study asking whether model value patterns merely imitate human norms or form distinct configurations. It supports the report's interest in structural comparison between model outputs and human baselines. Because this link is a mirror, use it to locate the canonical paper and version before citing results or extracting methods.

https://www.alphaxiv.org/abs/2404.12744

24 Evaluating Contextual Adaptation and Value Alignment in LLMs

PREPRINT

arXiv

A recent evaluation of whether value expression adapts appropriately to context. It is relevant to steerable and pluralistic alignment, where consistency cannot mean identical answers in every setting. As a preprint, its dataset construction, cultural coverage, evaluator model, and operational definition of 'appropriate' adaptation require close review.

https://arxiv.org/html/2609.38260v1

Computational values and AI alignment

The computational literature is growing quickly; peer-reviewed measurement papers deserve more weight than mirrors, indexes, or preprints.

25 Towards Pluralistic Alignment of LLMs: A Comprehensive Survey

PREPRINT SURVEY

Preprints.org

A survey organizing approaches to alignment under value diversity rather than a single universal target. It provides terminology and a map of methods useful for positioning Schwartz-based work within the broader field. Since the venue is a preprint server, verify the review's search strategy, coverage date, and later peer-reviewed version before treating it as comprehensive.

https://www.preprints.org/manuscript/202603.1876

26 Human Values and the Uneven Cultural Alignment of Artificial Intelligence

RESEARCH PAPER / REPOSITORY COPY

ResearchGate

A cross-cultural analysis of where AI value profiles align more closely with some populations than others. It directly supports the report's concern about Western, educated, industrialized, rich, and democratic bias. Check the canonical publication, sampling frame, languages, and country-level baselines before using it to compare specific products or regions.

https://www.researchgate.net/publication/415049598_Human_values_and_the_uneven_cultural_alignment_of_artificial_intelligence

27 From Values to Opinions: Predicting Human Behaviors and Stances from Value Preferences

PEER-REVIEWED CONFERENCE PAPER

EMNLP 2023

A natural-language processing study that links value distributions to downstream opinions and behavior. It is important for the report's discussion of value injection and population simulation. The paper offers a concrete operational pathway, but predictive success should not be mistaken for faithful representation of individuals, especially when values are inferred from limited text.

https://aclanthology.org/anthology-files/pdf/emnlp/2023.emnlp-main.961.pdf

Computational values and AI alignment

The computational literature is growing quickly; peer-reviewed measurement papers deserve more weight than mirrors, indexes, or preprints.

28 Assessing Language Models' Values Through Psychometrically Grounded Measures

PEER-REVIEWED CONFERENCE PAPER

ACL 2025

A measurement-focused study bringing psychometric discipline to language-model value assessment. It is central to moving beyond ad hoc prompt tests toward reliability, validity, and structural checks. This source should be prioritized when selecting instruments or interpreting differences between models, while still recognizing that self-report-style measures may diverge from open-ended behavior.

https://aclanthology.org/2025.acl-long.838.pdf

29 Measuring Human and AI Values Based on Generative Psychometrics

PEER-REVIEWED CONFERENCE PAPER

AAAI

A methodological paper on extracting value signals from free-form generations rather than forcing models into fixed-choice surveys. It underpins the report's account of Generative Psychometrics for Values. Its practical strength is ecological validity; its main audit points are classifier dependence, prompt coverage, aggregation rules, and whether automated labels reproduce the same biases they aim to measure.

https://ojs.aaai.org/index.php/AAAI/article/view/34839/36994

30 Awesome LLM in Social Science

CURATED RESOURCE LIST

GitHub / ValueByte-AI

A maintained discovery list covering language-model applications in social science. It is useful for finding datasets, methods, and neighboring literatures quickly. It is not a systematic review and may reflect maintainer selection, so use it for search expansion rather than evidence weighting.

https://github.com/ValueByte-AI/Awesome-LLM-in-Social-Science

Computational values and AI alignment

The computational literature is growing quickly; peer-reviewed measurement papers deserve more weight than mirrors, indexes, or preprints.

31 CLAVE: An Adaptive Framework for Evaluating Values of LLMs

PREPRINT

arXiv

A framework for adapting value evaluation to context and model behavior. It contributes a scalable alternative to rigid inventories and fits the report's argument that values are expressed through situational trade-offs. Before adoption, inspect how values are induced, who validates the labels, and whether adaptive evaluation improves construct validity rather than merely increasing coverage.

https://arxiv.org/html/2407.10725v1

32 Assessing the Alignment of Large Language Models With Human Values for Mental Health Integration

PEER-REVIEWED APPLIED STUDY

PMC

An applied evaluation of human-value alignment in a high-stakes domain. It shows how abstract value taxonomies become concrete safety and suitability questions. The domain focus is a strength and a limit: findings should guide mental-health-specific safeguards, not be generalized automatically to political, educational, or economic agents.

https://pmc.ncbi.nlm.nih.gov/articles/PMC11040439/

33 Large Language Model Psychometrics: A Systematic Review

PREPRINT SYSTEMATIC REVIEW

arXiv

A review of efforts to apply psychological instruments to language models. It provides a useful checklist of recurring validity problems, including prompt sensitivity, anthropomorphic interpretation, and uncertain comparability with human respondents. Its conclusions should be cross-checked against the final published version and the stated review protocol.

https://arxiv.org/html/2505.08245v3

Computational values and AI alignment

The computational literature is growing quickly; peer-reviewed measurement papers deserve more weight than mirrors, indexes, or preprints.

34 LLMs and Cultural Values: The Impact of Prompt Language and Context

PEER-REVIEWED JOURNAL ARTICLE

Computational Linguistics

A study of how language and context alter the cultural values expressed by models. It directly supports the report's warning that an English-centered model may not become culturally representative through translation alone. This is a high-priority source for multilingual benchmark design, especially for separating language effects from persona and population effects.

https://direct.mit.edu/coli/article/52/2/407/134313/LLMs-and-Cultural-Values-The-Impact-of-Prompt

35 Accepted Main Conference Papers - ACL 2024

CONFERENCE INDEX

ACL 2024

A program index for discovering relevant peer-reviewed conference papers. It helps verify venue placement and find adjacent work, but it is not an evidentiary source by itself. Cite the individual paper, not the program page, when making a substantive claim.

https://2024.aclweb.org/program/main_conference_papers/

36 Awesome LLM Psychometrics

CURATED RESOURCE LIST

GitHub / ValueByte-AI

A curated bibliography of language-model psychometrics research, tools, and datasets. It is valuable for keeping a fast-moving review current and for identifying measurement approaches beyond values. As with any 'awesome list,' coverage is selective and quality is heterogeneous; confirm primary sources and publication status independently.

https://github.com/ValueByte-AI/Awesome-LLM-Psychometrics

RESEARCH-USE MAP

Build the next review from the inside out

The report's central proposition is promising: Schwartz offers a structural vocabulary for value trade-offs that maps naturally onto computational evaluation. A defensible follow-on study should tighten the evidence chain at four points.

1 · Construct	Define whether the target is a model's stated preference, generated behavior, latent representation, or downstream effect on users. These are not interchangeable.
2 · Instrument	Use validated Schwartz measures where possible, preserve relative scoring, and test whether the expected circumplex emerges. Report reliability and prompt sensitivity.
3 · Comparison	Keep Schwartz values distinct from Moral Foundations Theory, personality, sentiment, and safety outcomes. Map relationships empirically rather than by label similarity.
4 · Generalization	Evaluate across languages, cultures, tasks, and interaction lengths. Separate steerability from stability, and alignment with a population from normative acceptability.

> **Practical evidence hierarchy:** primary theory and peer-reviewed measurement papers first; peer-reviewed applications second; well-specified preprints for frontier methods; repositories, indexes, and explainers for discovery only.

HIGHEST-VALUE VERIFICATION TASKS

- Confirm canonical publication records for entries hosted on ResearchGate, AlphaXiv, or author sites.
- Check whether 2026 preprints have newer versions or peer-reviewed successors.
- Reproduce benchmark scoring and structural tests before comparing named models.
- Document language, sampling, annotation, and evaluator-model choices.